

HIGH-PERFORMANCE Python

PART 2: SINGLE THREAD



High-performance Python

<https://github.com/squillero/python-lectures/>

Copyright © 2023 by Giovanni Squillero.

Permission to make digital or hard copies for personal or classroom use of these files, either with or without modification, is granted without fee provided that copies are not distributed for profit, and that copies preserve the copyright notice and the full reference to the source repository. To republish, to post on servers, or to redistribute to lists, contact the Author. These files are offered as-is, without any warranty.

september 2023

giovanni.squillero@polito.it

version 0.5

High-performance Python

3

Why HP Python?

- Python is not quite fast, at least not compared to C++ or Rust
- ... but it's not quite slow either
 - Large projects
 - System language
 - High-level algorithms
- Goal of HP Python
 - Avoid pitfalls
 - High-level optimization
 - Pure-Python optimization

giovanni.squillero@polito.it

High-performance Python



Rule of Optimization (personal)

- Don't write it
- Don't do it
- Exploit parallelism

giovanni.squillero@polito.it

High-performance Python



5

Rule of Optimization (personal)

- Don't write it
 - Use builtins
 - Use standard libraries
- Don't do it
- Exploit parallelism



giovanni.squillero@polito.it

High-performance Python

6

Rule of Optimization (personal)

- Don't write it
 - Use builtins
 - Use standard libraries
- Don't do it
 - Don't do it yet: i.e., use lazy execution
 - Don't do it again: i.e., cache previous calls
- Exploit parallelism



giovanni.squillero@polito.it

High-performance Python

7

Lazy Executions & Generators

```
def number_generator(n):
    i = 0
    while i < n:
        yield i
        i += 1

for number in number_generator(5):
    print(number)
```

giovanni.squillero@polito.it

High-performance Python

8

Problem: Write a Generator for Fibonacci



giovanni.squillero@polito.it

High-performance Python

9

Generators & Lazy Evaluation

- `any()`
- `all()`

giovanni.squillero@polito.it

High-performance Python

10

Problem: Calculate Prime Numbers

1	2	3	4	5	6	7	8	9	10
11	12	13	14	15	16	17	18	19	20
21	22	23	24	25	26	27	28	29	30
31	32	33	34	35	36	37	38	39	40
41	42	43	44	45	46	47	48	49	50
51	52	53	54	55	56	57	58	59	60
61	62	63	64	65	66	67	68	69	70
71	72	73	74	75	76	77	78	79	80
81	82	83	84	85	86	87	88	89	90
91	92	93	94	95	96	97	98	99	100

giovanni.squillero@polito.it

High-performance Python

11

DON'T DO IT

Out-of-the-box

New in
version 3.9

```
import functools

@functools.cache
def cached_function(num):
    print(f"Performing costly calculation for {num}")
    return num * 2

cached_function(42)
Performing costly calculation for 42
84

cached_function(42)
84

cached_function(43)
Performing costly calculation for 43
86
```

Out-of-the-box

```
@functools.lru_cache(maxsize=512)
def cached_function(num):
    print(f"Performing costly calculation for {num}")
    return num * 2

cached_function(42)
Performing costly calculation for 42
84

cached_function.cache_info()

CacheInfo(hits=0, misses=1, maxsize=128, currsize=1)

cached_function.cache_clear()

cached_function(42)
Performing costly calculation for 42
84
```

Out-of-the-box (caveats)

- **@functools.cached_property**
 - Similar to property, but with caching
- **@functools.lru_cache (typed)**
 - If **typed** is **True**, function arguments of different types will be cached separately
 - If **typed** is **False** (default), the implementation will “**usually**” regard different types as equivalent calls and only cache a single result — but some types “**may**” be always cached separately

Memoization: joblib

```
from joblib import Memory
memory = Memory('./cache', verbose=0)

@memory.cache
def memoized_function(num):
    print(f"Performing even more costly calculation for {num}")
    return num * 3

memoized_function(13)
Performing even more costly calculation for 13
39

memoized_function(13)
39
```

giovanni.squillero@polito.it

High-performance Python

16

Memoization: joblib

```
from joblib import Memory
memory = Memory('./cache', verbose=0)

@memory.cache
```

File/Folder	Type	Size
cache	File Folder	238 bytes
joblib	File Folder	
main-C%63A-Users-giova-AppData-Local-Temp-ipykernel-210146844	File Folder	
memoized_function	File Folder	
1e7c14b1db6bb925391eda230d02a8b	File Folder	
metadata.json	JSON File	92 bytes
output.pkl	PKL File	5 bytes
func_code.py	PY File	141 bytes

39

giovanni.squillero@polito.it

High-performance Python

17

Memoization: joblib

Memory (

```

    - location=None      // path/base directory
    - backend='local'   // local is regular file system
    - mmap_mode=None    // see numpy.memmap
    - compress=False     // compressed array cannot use memmap
    - verbose=1          // ...
    - bytes_limit=None  // int or string like "16M"
    - backend_options=None
)

```

giovanni.squillero@polito.it

High-performance Python

18



Memoization: joblib

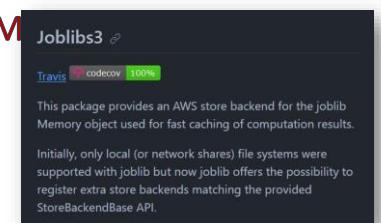
```

Memory(_STORE_BACKENDS = {'local': FileSystemStoreBackend})
    - location=None      // path/base directory
    - backend='local'   // local is regular file system
    - mmap_mode=None    // see numpy.memmap
    - compress=False     // compressed array cannot use memmap
    - verbose=1          // ...
    - bytes_limit=None  // int or string like "16M"
    - backend_options=None
)

```

giovanni.squillero@polito.it

High-performance Python



Memoization: joblib

```
from joblib import Memory
memory = Memory('~/cache', verbose=0)

@memory.cache(ignore=['foo'], verbose=2)
def memoized_function(num, foo):
    print(f"Performing even more costly calculation for {num}")
    return num * 3

memoized_function(13, foo=1)
```

[Memory] Calling __main__-C%3A-Users-giova-AppData-Local-Temp-ipykernel-499165028.memoized_function...
memoized_function(13, foo=1)
Performing even more costly calculation for 13
_memoized_function - 0.0s, 0.0min
39

```
memoized_function(13, foo=2)
```

[Memory] 776.5s, 12.9min : Loading memoized_function...
39

giovanni.squillero@polito.it

High-performance Python

20

Memoization: joblib

```
def my_function(*args, **kwargs):
    print(args, kwargs)
    return True

@memory.cache(cache_validation_callback=my_function, verbose=2)
def memoized_function(num):
    print(f"Performing even more costly calculation for {num}")
    return num * 3

memoized_function(13)
```

[Memory] Calling __main__-C%3A-Users-giova-AppData-Local-Temp-ipykernel-56899765.memoized_function...
memoized_function(13)
Performing even more costly calculation for 13
_memoized_function - 0.0s, 0.0min
39

```
memoized_function(13)
```

({'duration': 0.001997709274291992, 'input_args': {'num': '13'}, 'time': 1697901174.566357},)
[Memory] 41.1s, 0.7min : Loading memoized_function...
39

giovanni.squillero@polito.it

High-performance Python

21

Generic Persistence: pickle

```

import pickle

cache = dict()
def cached_function(num):
    if num not in cache:
        print("Performing costly calculation for %d" % num)
        cache[num] = num * 3 + 1
    return cache[num]

cached_function(42)
Performing costly calculation for 42
127

cached_function(42)
127

pickle.dump(cache, open('cache.pkl', 'wb'))

cache = pickle.load(open('cache.pkl', 'rb'))

cached_function(42)
127

```

giovanni.squillero@polito.it

High-performance Python

22

Generic Persistence: pickle

- Supports dumping to/loading from a string (dumps/loads)
- Supports different *protocols*
 - `pickle.DEFAULT_PROTOCOL`
 - 0 (zero) for almost human-readable data
 - `pickle.HIGHEST_PROTOCOL`

(dp0
I42
I127
S.

giovanni.squillero@polito.it

High-performance Python

23

Generic Persistence: pickle



giovanni.squillero@polito.it

High-performance Python

24

pickle vs. joblib

```
import joblib

joblib.dump(cache, open('cache.pkl', 'wb'))

cache = joblib.load(open('cache.pkl', 'rb'))

cached_function(42)
```

127

Note:

As of Python 3.8 and numpy 1.16, pickle protocol 5 introduced in [PEP 574](#) supports efficient serialization and de-serialization for large data buffers natively using the standard library:

```
pickle.dump(large_object, fileobj, protocol=5)
```

giovanni.squillero@polito.it

High-performance Python

25

pickle vs. marshal

- Python has a more primitive serialization module called **marshal**, but in general pickle should always be the preferred way to serialize Python objects
- **marshal** exists primarily to support Python's **.pyc** files

giovanni.squillero@polito.it

High-performance Python

26

Generic Persistence: shelve

```
import shelve

cache = None
def shelfed_function(num):
    key = str(num)
    if key not in cache:
        print(f"Performing costly calculation for {num}")
        cache[key] = num * 3 + 1
    return cache[key]

with shelve.open('cache.db') as cache:
    print(shelfed_function(42))
    print(shelfed_function(42))

Performing costly calculation for 42
127
127

with shelve.open('cache.db') as cache:
    print(shelfed_function(42))
    print(shelfed_function(42))

127
127
```

giovanni.squillero@polito.it

High-performance Python

27

Generic Persistence: `shelve`

```
import shelve

cache = None
def shelfed_function(num):
    key = str(num)
    if key not in cache:
        print(f"Performing costly calculation for {num}")
        cache[key] = num * 3 + 1
    return cache[key]

with shelve.open('cache.db') as cache:
    print(shelfed_function(42))
    print(shelfed_function(42))

Performing costly calculation for 42
127
127
```

The **keys** in a shelf are ordinary strings

The **values** can be anything that the pickle module can handle, including most class instances, recursive data types, and objects containing lots of shared sub-objects

```
with shelve.open('cache.db') as cache:
    print(cache['key'])
```

cache.db.bak	14 bytes	BAK File
cache.db.dat	5 bytes	DAT File
cache.db.dir	14 bytes	DIR File

giovanni.squillero@polito.it

High-performance Python

28

Bonus Topic: `weakref`

- A weak reference is not enough to keep an object alive
 - When the only remaining references to a referent are weak references, garbage collection may destroy the referent
- Note: Until the object is actually destroyed the weak reference may return the object

```
from dataclasses import dataclass
@dataclass
class Foo:
    num: int

x = Foo(23)
y = weakref.ref(x)

x, y, y()
```

(Foo(num=23), <weakref at 0x1067b9c60; to 'Foo' at 0x105f88390>, Foo(num=23))

giovanni.squillero@polito.it

High-performance Python

29

Bonus Topic: `weakref.proxy`

```
from dataclasses import dataclass
import weakref

@dataclass
class Foo:
    num: int

store = [Foo(n) for n in range(5)]
refs = [weakref.proxy(store[n]) for n in range(5)]

refs

[<weakproxy at 0x106070590 to Foo at 0x1068eb010>,
 <weakproxy at 0x106b5cc70 to Foo at 0x106076610>,
 <weakproxy at 0x106d50220 to Foo at 0x106077190>,
 <weakproxy at 0x106d513a0 to Foo at 0x106900fd0>,
 <weakproxy at 0x106d51530 to Foo at 0x106900a10>]

del store[2]
refs

[<weakproxy at 0x106070590 to Foo at 0x1068eb010>,
 <weakproxy at 0x106b5cc70 to Foo at 0x106076610>,
 <weakproxy at 0x106d50220 to NoneType at 0x103a65d48>,
 <weakproxy at 0x106d513a0 to Foo at 0x106900fd0>,
 <weakproxy at 0x106d51530 to Foo at 0x106900a10>]
```

giovanni.squillero@polito.it

High-performance Python

30

Bonus Topic: `weakref.proxy`

```
from dataclasses import dataclass
import weakref

@dataclass
class Foo:
    num: int

store = [Foo(n) for n in range(5)]
refs = [weakref.proxy(store[n]) for n in range(5)]

refs

[<weakproxy at 0x106070590 to Foo at 0x1068eb010>,
 <weakproxy at 0x106b5cc70 to Foo at 0x106076610>,
 <weakproxy at 0x106d50220 to Foo at 0x106077190>,
 <weakproxy at 0x106d513a0 to Foo at 0x106900fd0>,
 <weakproxy at 0x106d51530 to Foo at 0x106900a10>]

del store[2]
refs

[<weakproxy at 0x106070590 to Foo at 0x1068eb010>,
 <weakproxy at 0x106b5cc70 to Foo at 0x106076610>,
 <weakproxy at 0x106d50220 to NoneType at 0x103a65d48>,
 <weakproxy at 0x106d513a0 to Foo at 0x106900fd0>,
 <weakproxy at 0x106d51530 to Foo at 0x106900a10>]
```

giovanni.squillero@polito.it

High-performance Python

31

Bonus Topic: weak Dictionaries

- **WeakKeyDictionary**
 - Mapping class that references keys weakly
- **WeakValueDictionary**
 - Mapping class that references values weakly
- **keyrefs & valuerfs**
- **WeakSet**
 - Set class that keeps weak references to its elements

More Generic Persistence

- **dbm**
 - Interfaces to the traditional NoSQL “databases” available under Unix (i.e., fast, single-keyed, key-value store)
- **sqlite3**
 - DB-API 2.0 interface for SQLite



